


Feasibility Study:

FPGA Parallel Heterogeneous Computing

<https://github.com/yuwen41200/fpga-computing>



Why FPGAs ?

- FPGA = Field Programmable Gate Array
- FPGAs are essentially massively parallel processors
- Fully customized hardware
→ Maximized optimization
- Past research has also showed that using FPGAs as accelerators can realize better performance than using GPGPUs
- If the data transmission interface is also PCIe, the cost is almost the same as GPGPUs

Drawbacks

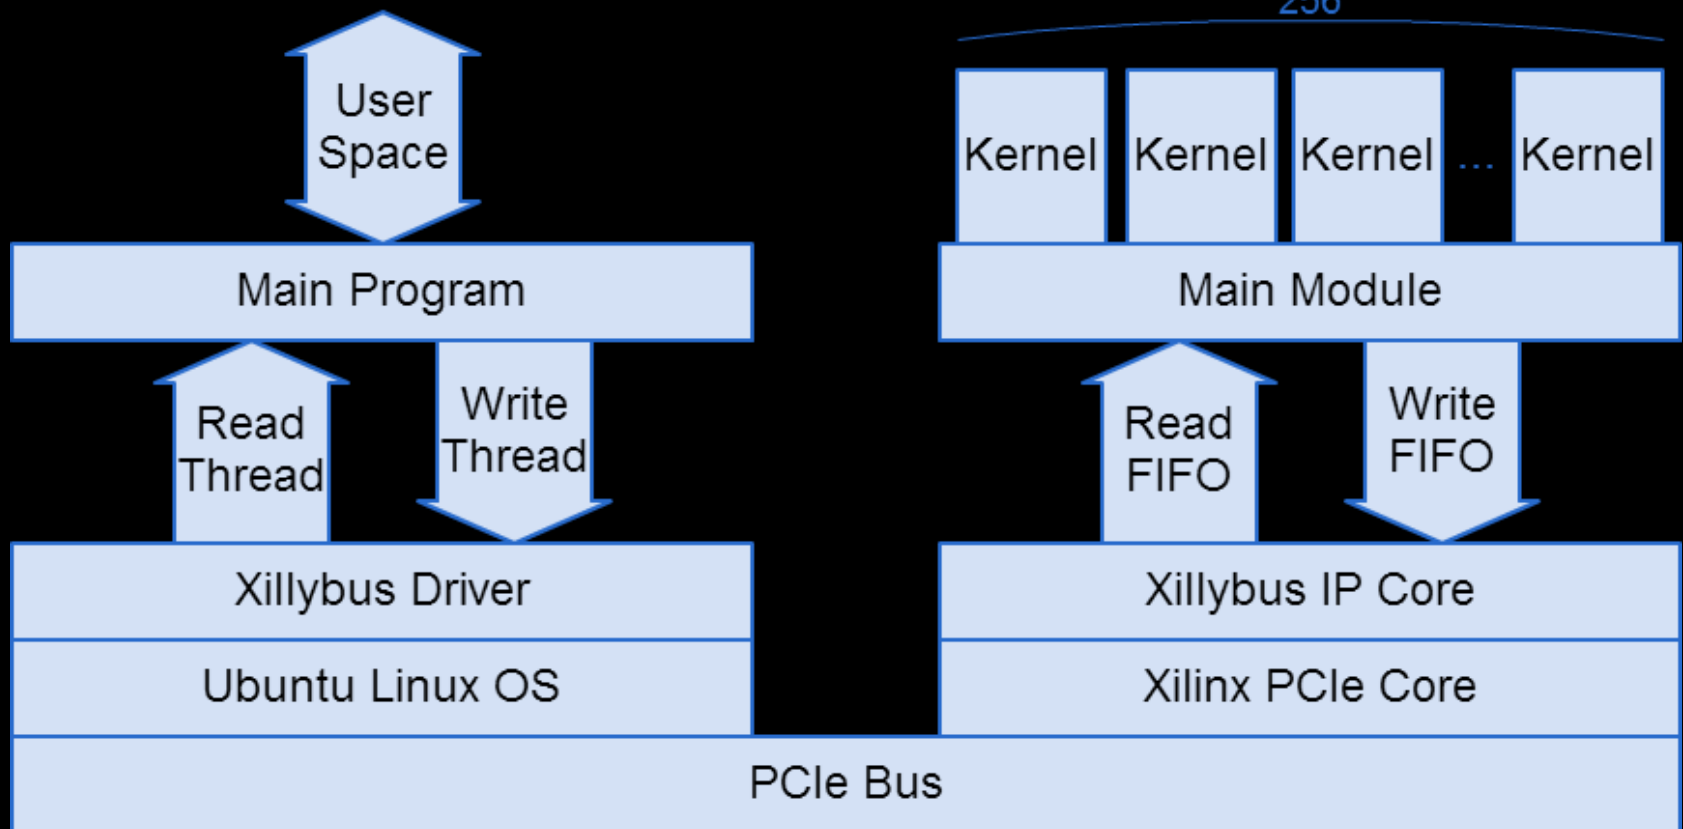
- Circuit design is hard, especially for experienced software developers
- The performance of the circuit is highly dependent on the quality of the circuit design and the available resources on the FPGA development board
- Though some EDA (electronic design automation) tools, like Xilinx SDAccel, support high-level synthesis that can convert OpenCL C/C++ codes into schematics, these techniques are not mature yet, and may lead to poor performance

Proposed Implementation

- Employ RTL (register-transfer level) design by the Verilog HDL (hardware description language)
- Use the Xillybus IP core for data transmission over the PCIe interface
- Targeted board: Xilinx Virtex-5 ML506 Evaluation Platform
- Use C++ for software design
- The host programs run on 64-bit Linux distributions

PC

FPGA



Problem Statement

```
14 always @(posedge clk) begin
15     if (in_valid && counter == 0) begin
16         result <= in_data;
17         counter <= counter + 1;
18         out_valid <= 0;
19     end
20     else if (in_valid && counter < 1025) begin
21         result <= result + (result >> 3);
22         counter <= counter + 1;
23         out_valid <= 0;
24     end
25     else if (in_valid && counter == 1025) begin
26         out_valid <= 1;
27     end
28     else if (~in_valid) begin
29         counter <= 0;
30         out_valid <= 0;
31     end
32 end
```

$$color \times 1.125^{1024}$$

← Hardware Version
↓ Software Version

```
30     for (int i = 0; i < 1920; ++i)
31         for (int j = 0; j < 1080; ++j)
32             for (int k = 0; k < 1024; ++k)
33                 frame0[i][j] = frame0[i][j] + (frame0[i][j] >> 3);
```

Kernel Instantiation

```
283 generate
284     genvar thread_no;
285     for (thread_no = 0; thread_no < THREAD_NUMBER; thread_no = thread_no + 1) begin:warp
286         kernel kernel_ins (
287             .clk(bus_clk),
288             .in_data(in_data[thread_no]),
289             .in_valid(in_valid[thread_no]),
290             .out_data(out_data[thread_no]),
291             .out_valid(out_valid[thread_no])
292         );
293     end
294 endgenerate
```

FSM (Finite-State Machine)

```
191  always @(*) begin
192      case (curr_state)
193          IDLE_STATE:
194              if (user_w_write_32_open && user_r_read_32_open)
195                  next_state = RECV_STATE;
196              else
197                  next_state = IDLE_STATE;
198          RECV_STATE:
199              if (recv_counter == THREAD_NUMBER-2)
200                  next_state = EXEC_STATE;
201              else
202                  next_state = RECV_STATE;
203          EXEC_STATE:
204              if (exec_done)
205                  next_state = SEND_STATE;
206              else
207                  next_state = EXEC_STATE;
208          SEND_STATE:
209              if (send_counter == THREAD_NUMBER-2)
210                  next_state = IDLE_STATE;
211              else
212                  next_state = SEND_STATE;
213      endcase
214  end
```


Data Interface

```
9  const char *readDeviceFile = "/dev/xillybus_read_32";
10 const char *writeDeviceFile = "/dev/xillybus_write_32";
11
12 int initRead() {
13     int fileDescriptor = open(readDeviceFile, O_RDONLY);
14     if (fileDescriptor < 0) {
15         if (errno == ENODEV)
16             fprintf(stderr, "Maybe %s is a write-only file.\n", readDeviceFile);
17         perror("Failed to open the device file for read");
18         exit(1);
19     }
20     return fileDescriptor;
21 }
```

Multithreading on Software

```
21 void sender() {
22     int fd = initWrite();
23     for (int i = 0; i < 1920; ++i)
24         fpgaWrite(fd, (unsigned char*) frame0[i], 2160);
25     for (int i = 0; i < 1920; ++i)
26         fpgaWrite(fd, (unsigned char*) frame1[i], 2160);
27     for (int i = 0; i < 1920; ++i)
28         fpgaWrite(fd, (unsigned char*) frame2[i], 2160);
29 }
30
31 void receiver() {
32     int fd = initRead();
33     for (int i = 0; i < 1920; ++i)
34         fpgaRead(fd, (unsigned char*) frame3[i], 2160);
35     for (int i = 0; i < 1920; ++i)
36         fpgaRead(fd, (unsigned char*) frame4[i], 2160);
37     for (int i = 0; i < 1920; ++i)
38         fpgaRead(fd, (unsigned char*) frame5[i], 2160);
39 }
```

Platform Specification

- **Serial:** Intel Core i5-3570 CPU @ 3.40GHz 4C4T, 8 GB Memory, Ubuntu 14.04 LTS 64-bit
- **CPU:** Intel Core i5-3570 CPU @ 3.40GHz 4C4T, 8 GB Memory, Ubuntu 14.04 LTS 64-bit, using up to 4 threads
- **GPGPU:** Intel Core i7-3770 CPU @ 3.40Ghz 4C8T, 16 GB Memory, Windows 10 Enterprise 64-bit, NVIDIA GeForce GTX 670
- **FPGA:** Intel Core i5-3570 CPU @ 3.40GHz 4C4T, 8 GB Memory, Ubuntu 14.04 LTS 64-bit, Xilinx Virtex-5 ML506 Evaluation Platform

Results

Data Size	Serial (Singlethread)	CPU Multithreading (OpenMP)	GPGPU Accelerator (CUDA)	FPGA Accelerator
Small (128 KiB)	0.179096 s	0.067667 s	0.891333 s	0.000534 s
<i>Speedup</i>	1.000 x	2.647 x	0.201 x	335.386 x
Large (~12 MiB)	15.873200 s	4.678667 s	1.799333 s	0.016104 s
<i>Speedup</i>	1.000 x	3.393 x	8.822 x	985.668 x

Attribution

Special thanks to Chun-Jen Tsai,
assistant professor at Dept. of C.S., NCTU

